



**National University**  
Of Computer and Emerging Sciences

# **PROJECT**

## **TORCS MACHINE LEARNING PROJECT**

In partial fulfillment of the requirement  
for the course of

***ARTIFICIAL INTELLIGENCE  
(AI-2002)***

By

**M. Huzaifa (22i-  
1220)**

**M. Sarmad (22i-0941)**

**M. Umar Hassan  
(22i-0942)**

**(CS-H)**

## Table of Contents

<b>1. Data Pre-processing .....</b>	<b>3</b>
<b>1.1 Loading &amp; Filtering .....</b>	<b>3</b>
<b>1.2 Dropping Unused Columns .....</b>	<b>3</b>
<b>1.3 Defining Features &amp; Labels.....</b>	<b>3</b>
<b>1.4 Normalization .....</b>	<b>4</b>
<b>1.5 Code:.....</b>	<b>5</b>
<b>2. Model Training .....</b>	<b>5</b>
<b>2.1 Code:.....</b>	<b>7</b>
<b>3. Inference &amp; Car Control .....</b>	<b>7</b>
<b>3.1 Load Artifacts .....</b>	<b>7</b>
<b>3.2 Feature Vector Construction .....</b>	<b>7</b>
<b>3.3 Normalization .....</b>	<b>8</b>
<b>3.4 Prediction .....</b>	<b>8</b>
<b>3.5 Control Mapping.....</b>	<b>8</b>
<b>3.6 Feedback Loop .....</b>	<b>8</b>
<b>4. Work Division:.....</b>	<b>8</b>

# TORCS Machine Learning Project

## Abstract:

In this project, we developed an intelligent, data-driven car racing controller using the TORCS framework. By transforming raw telemetry data into clean, normalized feature–label pairs, we enabled supervised learning for continuous control prediction using Multilayer Perceptron (MLP). The use of modern training practices—such as validation splits, normalization, and regularization callbacks, helped build a robust model that can effectively navigate various tracks under competitive conditions.

## 1. Data Pre-processing

In this project we transform raw TORCS telemetry logs into normalized input–output pairs for supervised learning. Key steps:

### 1.1 Loading & Filtering

- Read Dataset.csv into a pandas DataFrame.
- Remove “pre-start” rows where CurrentLapTime < 0.

### 1.2 Dropping Unused Columns

Discard meta-columns not directly useful for driving control.

Damage, CurrentLapTime, DistanceFromStart, DistanceCovered, FuelLevel, LastLapTime, RacePosition.

### 1.3 Defining Features & Labels

- **Features (X):**
  - Car orientation & position: Angle, TrackPosition, Z
  - Speeds: SpeedX, SpeedY, SpeedZ, RPM
  - Opponents: Opponent\_1 ... Opponent\_36
  - Track sensors: Track\_1 ... Track\_19

- Wheel spin: WheelSpinVelocity\_1 ... WheelSpinVelocity\_4
- Gear (included as input)
- **Labels (y):**
- Continuous control targets: Clutch, Braking, Steering, Acceleration, Gear

## 1.4 Normalization

- Fit a MinMaxScaler on all features to rescale into [0,1].
- Save the fitted scaler (scaler\_input.pkl) for inference.
- **Train/Validation Split**
- Perform an 80/20 split (random seed 42).
- Export CSV files: train\_X.csv, val\_X.csv, train\_y.csv, val\_y.csv.

```

Loading raw data from: ./Dataset.csv
Dropped pre-start rows; remaining rows: 912979
Fitted MinMaxScaler including gear and saved scaler.
Preprocessing complete:
train_X shape: (730383, 66)
val_X shape:   (182596, 66)
train_y shape: (730383, 4)
val_y shape:   (182596, 4)

```

*Figure 1 Data pre-processing*

## 1.5 Code:

```
11 # 1) Load raw TORCS log data
12 print("Loading raw data from:", RAW_CSV)
13 df = pd.read_csv(RAW_CSV)
14
15 # 2) Drop countdown / pre-start rows (curLapTime < 0)
16 df = df[df["CurrentLapTime"] >= 0].reset_index(drop=True)
17 print("Dropped pre-start rows; remaining rows:", df.shape[0])
18
19 # 3) Drop unused / meta columns (keep 'gear' as feature)
20 drop_cols = [
21     "Damage", "CurrentLapTime", "DistanceFromStart",
22     "DistanceCovered", "FuelLevel", "LastLapTime",
23     "RacePosition"
24 ]
25 df = df.drop(columns=drop_cols)
26
27 # 4) Define feature & label columns, include 'gear' as feature
28 track_cols = [f"Track_{i}" for i in range(1,20)]
29 wsv_cols = [f"WheelSpinVelocity_{i}" for i in range(1,5)]
30 opp_cols = [f"Opponent_{i}" for i in range(1,37)]
31
32 feature_cols = ["Angle"] + opp_cols + [
33     "RPM", "SpeedX", "SpeedY", "SpeedZ", "TrackPosition", "Z"
34 ] + track_cols + wsv_cols
35 label_cols = ["Clutch", "Braking", "Steering", "Acceleration"] # gear removed from label
36
37 # 5) Normalize features into [0,1] using MinMaxScaler
38 scaler = MinMaxScaler()
39 X_norm = scaler.fit_transform(df[feature_cols])
40 joblib.dump(scaler, SCALER_PATH)
41 print("Fitted MinMaxScaler including gear and saved scaler.")
42
43 # 6) Prepare labels (gear is now input, so labels exclude it)
44 y_df = df[label_cols].reset_index(drop=True)
45
```

---

## 2. Model Training

**Algo Chosen :** Neural Networks

**Reason :** Because we had multiple labels to infer  
and also the dependencies were complex

Trained a deep feed-forward neural network to map the N-dimensional feature vector to the 5dimensional control vector:

- **Network Architecture**

Sequential Keras model with nine Dense layers:

Input → 1024 → 512 → 256 → 128 → 64 → 32 → 16 → 8 → Output(5) Hidden  
layers use ReLU activations; the output layer is linear.

- **Compilation Settings**

- Optimizer: Adam (**learning rate = 1e-4**)
- Loss: Mean Squared Error (MSE)
- Metric: Mean Absolute Error (MAE)
- **Callbacks**
  - ModelCheckpoint: save best model by validation loss
  - EarlyStopping: stop if no improvement for 5 epochs, restore best weights
  - ReduceLROnPlateau: halve LR on plateau (factor 0.5, patience 3, min\_lr=1e6)
- **Training Run**
  - Epochs: up to 50
  - Batch size: 64
  - Validate on 20% hold-out each epoch
  - Best model saved to trained\_model.h5

```

11413/11413 — 267s 23ms/step - loss: 0.0021 - mae: 0.0179 - val_loss: 0.0034 - val_mae: 0.0204 - learning_rate: 6.2500e-06
Epoch 42/50
11412/11413 — 0s 23ms/step - loss: 0.0020 - mae: 0.0178
Epoch 42: ReduceLROnPlateau reducing learning rate to 3.12499992105586e-06.
11413/11413 — 273s 24ms/step - loss: 0.0020 - mae: 0.0178 - val_loss: 0.0034 - val_mae: 0.0203 - learning_rate: 6.2500e-06
Epoch 43/50
11413/11413 — 0s 23ms/step - loss: 0.0020 - mae: 0.0175
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend
11413/11413 — 1176s 103ms/step - loss: 0.0020 - mae: 0.0175 - val_loss: 0.0033 - val_mae: 0.0201 - learning_rate: 3.1250e-06
Epoch 44/50
11413/11413 — 968s 85ms/step - loss: 0.0020 - mae: 0.0175 - val_loss: 0.0033 - val_mae: 0.0201 - learning_rate: 3.1250e-06
Epoch 45/50
11412/11413 — 0s 22ms/step - loss: 0.0020 - mae: 0.0175
Epoch 45: ReduceLROnPlateau reducing learning rate to 1.56249996052793e-06.
11413/11413 — 271s 24ms/step - loss: 0.0020 - mae: 0.0175 - val_loss: 0.0033 - val_mae: 0.0201 - learning_rate: 3.1250e-06
Epoch 46/50
11411/11413 — 0s 23ms/step - loss: 0.0019 - mae: 0.0173
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend
11413/11413 — 278s 24ms/step - loss: 0.0019 - mae: 0.0173 - val_loss: 0.0033 - val_mae: 0.0200 - learning_rate: 1.5625e-06
Epoch 47/50
11413/11413 — 264s 23ms/step - loss: 0.0019 - mae: 0.0173 - val_loss: 0.0033 - val_mae: 0.0200 - learning_rate: 1.5625e-06
Epoch 48/50
11411/11413 — 0s 22ms/step - loss: 0.0019 - mae: 0.0173
Epoch 48: ReduceLROnPlateau reducing learning rate to 1e-06.
11413/11413 — 271s 24ms/step - loss: 0.0019 - mae: 0.0173 - val_loss: 0.0033 - val_mae: 0.0200 - learning_rate: 1.5625e-06
Epoch 49/50
11411/11413 — 0s 23ms/step - loss: 0.0019 - mae: 0.0172
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend
11413/11413 — 279s 24ms/step - loss: 0.0019 - mae: 0.0172 - val_loss: 0.0033 - val_mae: 0.0200 - learning_rate: 1.0000e-06
Epoch 50/50
11413/11413 — 254s 22ms/step - loss: 0.0019 - mae: 0.0172 - val_loss: 0.0033 - val_mae: 0.0200 - learning_rate: 1.0000e-06
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend
Model training complete. Best model saved to ./trained_model.h5

```

Figure 2: Model Training

## 2.1 Code:

```
22 model = Sequential([
23     Dense(1024, activation='relu', input_shape=(X_train.shape[1],)),
24     Dense(512, activation='relu'),
25     Dense(256, activation='relu'),
26     Dense(128, activation='relu'),
27     Dense(64, activation='relu'),
28     Dense(32, activation='relu'),
29     Dense(16, activation='relu'),
30     Dense(8, activation='relu'),
31     Dense(y_train.shape[1])
32 ])
33
34 # 3) Compile model with learning rate 1e-4
35 model.compile(
36     optimizer=Adam(learning_rate=1e-4),
37     loss='mse',
38     metrics=['mae']
39 )
40
41 # 4) Set up callbacks including LR reduction on plateau
42 callbacks = [
43     ModelCheckpoint(MODEL_PATH, save_best_only=True, monitor='val_loss'),
44     EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True),
45     ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6, verbose=1)
46 ]
47
48 # 5) Train model
49 history = model.fit(
50     X_train, y_train,
51     validation_data=(X_val, y_val),
52     epochs=50,
53     batch_size=64,
54     callbacks=callbacks
55 )
```

---

## 3. Inference & Car Control

At runtime, the trained model and scaler are used to compute control commands from live sensor data:

### 3.1 Load Artifacts

- Load MinMaxScaler to replicate training-time normalization.
- Load the Keras model (trained\_model.h5) in inference mode.

### 3.2 Feature Vector Construction

- Read current state: angle, opponent distances, speeds, track sensors, wheel spins, gear.

- Assemble into a 1×N NumPy array.

### 3.3 Normalization

- $X\_norm = (X\_raw - scaler.data\_min\_)/ (scaler.data\_max\_ - scaler.data\_min\_)$

### 3.4 Prediction

- Model outputs: [clutch, brake, steering, acceleration, gear].

### 3.5 Control Mapping

- Clamp each value to valid range (e.g., steering  $\in [-1,1]$ , accel  $\in [0,1]$ ).
- Send commands to the car control interface.

### 3.6 Feedback Loop

- Repeat for each incoming sensor message to form a closed-loop controller

## 4. Work Division:

- **Data Collection** : Muhammad Huzaifa
- **Pre Processing** : Muhammad Umar Hassan
- **Model Training & Driver Implementation** : Muhammad Sarmad